

P \neq NP

Sheldon S. Nicholl

May 21, 2025

Abstract

Not for publication! Still in draft form; probably contains errors.

In this paper, finite state automata and P-Complete problems are used to study the P versus NP problem. First, all the states at a particular depth of the minimal automaton for a restricted Clique problem are shown to have an in-degree of 1. Next, any automaton for any P-Complete language for strings of the same length is shown to have at least one state at that same level with in-degree greater than 1. So there can be no such circuit and P does not equal NP.

1 Definitions

An *alphabet* is a finite set of characters. A *string* is a finite sequence of characters drawn from an alphabet. A *language* is a set of strings.

A *bit string* is an element of the language $\{0, 1\}^*$. If a string $s = vw$, then v is a *prefix* of s and w is a *suffix* of s . The i th character c of string s is accessed with the following notation: $c = s[i]$.

For a language L , define $L_n \subseteq L$ to consist of strings in L of length n :

$$L_n = \{s \mid s \in L \text{ and } |s| = n\} \quad (1)$$

The function v takes a graph g and returns its vertices, so if $g = (V, E)$ then $v(g) = V$. The notation $g_1 \cong g_2$ is used to indicate that graphs g_1 and g_2 are isomorphic.

Following Sipser [2], a deterministic finite automaton (DFA) is a tuple

$$A = (\Sigma, S, F, \delta, q_0) \quad (2)$$

where Σ is a finite alphabet, S is a finite set of states, $F \subseteq S$ is the set of final states, $\delta : S \times \Sigma \rightarrow S$ is a mapping called the transition function, and $q_0 \in S$ is the initial state.

A state q is considered *accessible* if there is a path from the start state to q . A state q is called *co-accessible* if there is a path from q to a final state. Finally, an automaton is called *trim* if all its states are both accessible and co-accessible. To accommodate trim automata, the state transition function δ is allowed to be a partial function. All automata in this paper are presumed to be minimal, trim, and acyclic unless otherwise specified. By contrast, an automaton with a total transition function will be called a *total automaton*.

If $a \in \Sigma$ and the state $s' = \delta(s, a)$ exists, then the pair (s, s') can be viewed as an edge in a graph. So the automaton defines a directed graph with vertices S and edges based on δ . The usual graph concepts like a path, shortest path, and depth then come in their usual meanings. The *depth* of a state $s \in S$ is the length of the shortest path from the start state q_0 to s . The function $d : S \rightarrow \mathbb{N}$ will give the numeric depth of a given state. So for example, $d(q_0) = 0$. A *level* is the set of all states in S at a given depth. The notation $|M|$ refers to the *size* of the automaton: the number of states in the automaton.

A Boolean circuit and a finite-state automaton will be called *equivalent* if they decide the same language of bit strings.

2 Automata to Decide the k-Clique Problem

Define $\Gamma(n, k)$ as the set of graphs of size n that have a clique of size k :

$$\Gamma(n, k) = \{g \mid g \text{ contains a clique of size } k, |g| = n\} \quad (3)$$

Define $\Gamma(n)$ to be the set of graphs of size n that have a clique of size $\lfloor \frac{n}{2} \rfloor - 1$:

$$\Gamma(n) = \Gamma(n, \lfloor \frac{n}{2} \rfloor - 1) \quad (4)$$

The set $\Gamma(n)$ is large:

$$|\Gamma(n)| = \Omega(2^n) \quad (5)$$


2.1 Automaton construction

A finite state automaton can be created from a finite language via the following standard procedure. Since the language is finite, it can be written $\{w_1, w_2, w_3, \dots, w_m\}$ where $w_i \in \Sigma^*$. This set can then be turned into the following regular expression: $E = w_1 \cup w_2 \cup w_3 \cup \dots \cup w_m$. E is then converted to a non-deterministic finite automaton N . Then N is converted into a deterministic finite automaton D . Finally, D is minimized to yield the minimal automaton [2].

A finite-state automaton to decide the Clique problem $\Gamma(n, k)$ will now be constructed. The entire graph $g \in \Gamma(n, k)$ will be encoded as a string. An adjacency matrix representation will be used to encode g over the alphabet $\{0, 1\}$, i.e., a bit string.

2.2 Encoding a graph

The encoding of graphs is described below in some detail. An adjacency matrix will be used to encode a graph $g \in \Gamma(n, k)$. The following graph only intended to illustrate the encoding process and does not capture the $2n$ argument being made here:



Since the graph is undirected, the adjacency matrix is symmetric, making the top and bottom triangular matrices redundant, so only the upper triangular matrix will be encoded, which looks like this:

	1	1	0	0	r_1
		1	1	0	r_2
			0	1	r_3
				0	r_4
					r_5

Row-major order will be used to flatten the matrix into a linear form. Consider a column c_i . Each digit in c_i will be treated as a character and all those characters will be concatenated together to form a string s_i . So in this example, an empty string shows up since column 1 is on the diagonal:

$$s_1 = 1100 \quad (6)$$

$$s_2 = 110 \quad (7)$$

$$s_3 = 01 \quad (8)$$

$$s_4 = 0 \quad (9)$$

$$s_5 = \epsilon \quad (10)$$

Finally all the s_i are concatenated to form the encoding θ :

$$\theta = s_1 \cdot s_2 \cdot s_3 \cdot s_4 \cdot s_5 = 1100 \cdot 110 \cdot 01 \cdot 0 \cdot \epsilon = 1100110010 \quad (11)$$


Now define a function $e : V \times E \rightarrow \Sigma^*$ which maps graphs into their encodings as described above, so for example, $e(g) = \theta$. The function e extends naturally to sets of graphs, so applying e to a set of graphs G yields the set of encodings: $e(G) = \{e(g) \mid g \in G\}$.

Now let $\Theta_{n,k}$ be the encodings of those graphs for $\Gamma(n, k)$:

$$\Theta_{n,k} = \{\theta \mid g \in \Gamma(n, k) \text{ and } e(g) = \theta\}$$

2.3 Automata to decide an encoding

There is an automaton to decide $\Theta_{n,k}$. Now since $\Theta_{n,k}$ is a finite set of finite strings, it follows from the construction above that there exists a trim acyclic deterministic finite-state automaton $M_{n,k} = (\Sigma, S, q_0, \delta, F)$ to decide it: $M_{n,k}$ accepts θ iff $g \in \Gamma(n, k)$. In the running example, the automaton looks like this:



2.4 Lemma




Figure 1: Graph g with two 3-vertex structures in a clique of size 6




Figure 2: Graph g' , like graph g , but with edge (a,b) removed

Lemma 2.1. *All the states in the minimal automaton for $\Gamma(n, k)$ at level $k-1$ for $k < \lfloor \frac{n}{2} \rfloor$ in $M_{n,k}$ have in-degree 1, and so do all their ancestor states.*

Proof. Consider a graph $g \in \Gamma(n, k)$, where $g = (V, E)$. Let the vertices in V be numbered, so V can be written $V = \{v_1, v_2, v_3, \dots, v_n\}$. While there may be multiple cliques in a graph, it is also possible for there to be only one. So suppose that there exists a single clique $C \subseteq V$ where $|C| = k$ and $k < \lfloor \frac{n}{2} \rfloor$. Let $m = k - 1$. Let $g_t = (V_t, E_t)$ be the vertex-induced subgraph of g such that $V_t = \{v_1, v_2, v_3, \dots, v_m\}$; see Figure 1 for a visualization of this. Assume some of the vertices in C occur among the first m vertices in V . Let the string encoding of g_t be s_t , so $e(g_t) = s_t$. The string encoding of the whole graph g is by definition $e(g)$. Now define s_b as that suffix of $e(g)$ such that $e(g) = s_t \cdot s_b$.

Now consider a graph $g' = (V, E')$ which is equal to g in all respects except an edge from v_p to v_q in C has been removed, where $p \leq m$ and $q \leq m$ and $p < q$; see Figure 2 for a visualization of this. Let $g'_t = (V'_t, E'_t)$ be the vertex-induced subgraph of g' such that $V'_t = \{v'_1, v'_2, v'_3, \dots, v'_m\}$. So the clique is gone. This means the corresponding subgraph g'_t has changed, but the rest of g' has not. Let $e(g'_t) = s'_t$. Since $g_t \neq g'_t$, it follows that $e(g_t) \neq e(g'_t)$ so $s_t \neq s'_t$. Now define s'_b as that suffix of $e(g')$ such that $e(g') = s'_t \cdot s'_b$. But since nothing changed from g to g' outside of g_t , this means that $s_b = s'_b$. So $e(g') = s'_t \cdot s_b$.

This is the setup for the Myhill-Nerode argument that follows. There is the membership $s_t \cdot s_b \in \Theta_{n,k}$ since there is a clique in g . But $s'_t \cdot s_b \notin \Theta_{n,k}$ since there is no clique in g' . These two membership relations show that s_b is a Myhill-Nerode distinguishing extension for s_t and s'_t . So s_t and s'_t cannot be in the same Myhill-Nerode equivalence class. Therefore by the Myhill-Nerode theorem, they cannot be in the same state in a minimal finite-state machine to recognize $\Theta_{n,k}$. So all the states at level $k - 1$ are pairwise distinct, and therefore their in-degree must be 1. The ancestor states must also all be distinct since otherwise there would still be pairs of indistinguishable states at level $k - 1$. \square

Remark. Some details of the prior Lemma warrant explanation. The size of k is bounded by $\lfloor \frac{n}{2} \rfloor$ to ensure there are enough Myhill-Nerode distinguishers to distinguish all the prefixes. The expression $m = k - 1$ is introduced because the first row of the adjacency matrix could form the prefix encoding of a clique of size $k + 1$, hence the correction.

3 Relation to P-Complete languages

Consider a family $\mathcal{L} = \{L_1, L_2, L_3, \dots\}$ of languages where the subscripts describe the length of the strings in the languages as defined in Equation 1. Assume that $\cup L_n$ is not a sparse language and that all the L_n are decidable in polynomial time. So by definition, there is at least a polynomial-time reduction from any language in \mathcal{L} to a P-complete language like the Circuit Value Problem, where P_1, P_2, P_3, \dots is a family of Boolean circuits. Each Boolean circuit P_n has only a polynomial number of gates as a function of n : $|P_n| = O(n^k)$ for a constant k .

There are reductions from Circuit Value to context-free grammar recognition and unit resolution, so most of the reasoning here will be based on these two problems.

Definition 3.1. A context-free grammar is a quadruple $G = (V, \Sigma, R, S)$ following the usual definitions [2]. The set V is the non-terminals, and Σ is the set of terminals.

Definition 3.2. A divergent context-free language family is a family

$$\mathcal{L} = \{L_1, L_2, L_3, \dots\}$$

of context-free languages over bit strings where (1) $\cup L_n$ is not sparse, and (2) each L_n is decidable by a context-free grammar G_n that is polynomially bounded, i.e., $|G_n| = O(n^k)$.

Lemma 3.1. *Given a context-free grammar $G_n = (V_n, \Sigma_n, R_n, S_n)$ which decides a language L_n of $\Omega(2^n)$ binary strings where $|V_n| = O(n^k)$ for fixed k , it follows that a string of non-terminals $v_n \in V_n^*$ cannot encode an arbitrary n -bit string from L_n if $|v_n| = o(n/\log(n))$. Otherwise at least two of the n -bit strings will have the same encoding.*

Proof. Since $|L_n| = \Omega(2^n)$, it requires $\Omega(n)$ bits of information to distinguish all the strings in L_n correctly. But since $|V_n| = O(n^k)$, only $O(\log(n))$ bits of information come from each element of V_n . So by the Pigeonhole principle, as n gets large enough, there will be a situation where multiple bit strings will map to a single sequence of non-terminals since $|v_n| = o(n/\log(n))$. \square

Now consider a set of clauses to decide L_n using unit resolution. The inputs will be encoded with unit clauses. Since the alphabet of L_n is $\{0, 1\}$, a clause for 1 can be positive while a clause for 0 will be negative. A string in the language to be recognized just consists of ones and zeros, which will be encoded into the unit resolution formalism as a particular sequence of unit clauses.

Lemma 3.2. *Given a set of clauses C_n which decides a language L_n of $\Omega(2^n)$ binary strings with $\Omega(2^{\frac{n}{2}})$ prefix strings of size $\frac{n}{2}$ using unit resolution where $|C_n| = O(n^k)$ for fixed k , a set of resolvents C'_n cannot encode an arbitrary $\frac{n}{2}$ -bit binary prefix string from L_n if $|C'_n| = o(n/\log(n))$. Otherwise at least two of the $\frac{n}{2}$ -bit strings will receive the same encoding. Since $\lfloor \frac{n}{2} \rfloor - 1 = O(\frac{n}{2})$, this result applies to prefix strings of size $\lfloor \frac{n}{2} \rfloor - 1$ as well.*

Proof. By Lemma 3.1 and the reduction from the unit resolution problem to context-free grammar recognition ([1]). \square

Lemma 3.3. *Let $m = \lfloor \frac{n}{2} \rfloor - 1$. Given a language L_n in a divergent context-free language family and n sufficiently large, the finite-state automaton to decide L_n must always have a state at level m with in-degree greater than 1.*

Proof. Since L_n is finite, it can be recognized by a finite-state automaton, but the construction here will model the process of unit resolution on a finite language. Equation 2 showed the definition of a deterministic finite-state automaton, which will be repeated here: $A = (\Sigma, S, F, \delta, q_0)$. The alphabet is binary since L_n is in a divergent context-free language family, so $\Sigma = \{0, 1\}$. A state is a set of clauses to be defined shortly. If I is the set of input unit clauses and C_n is the set of clauses to decide language membership, then the start state q_0 then is just $I \cup C_n$. The final state is just the empty clause, so the set of final states F is just a set containing the empty clause: $F = \{\{\}\}$. A transition in the transition function δ will model resolution against a unit clause as follows. The start state q_s of the transition will be a set of clauses. The character will be 1 or 0 mapped to a unit clause c . The target state q_t of the transition will be that set of clauses arising from the application of unit propagation to the clauses in q_s with the unit clause c , so the transition is $q_s \xrightarrow{c} q_t$. The transition function consists of all such transitions starting from the start state. Finally, the set of all states S will be the transitive closure of all the possible transitions from the start state.

So by Lemma 3.2 with n sufficiently large, there must be a set of clauses R such that multiple inputs of length m land there. But since R is just a state in the automaton construction of the previous paragraph, this state, which is also at level m , must have in-degree greater than 1. \square

4 Conclusions

Theorem 4.1. $P \neq NP$.

Proof. Let $m = \lfloor \frac{n}{2} \rfloor - 1$. Proof is by reductio ad absurdum. So assume there is a polynomial-time algorithm for Clique. So by this assumption there is a circuit family

$$C = \{C_1, C_2, C_3, \dots, C_n, \dots\}$$

to decide the Clique problem as defined in Equation 4 which grows as $|C_n| = O(n^c)$ for some fixed c . Under this assumption and the reduction from Circuit Value to Context-Free Grammar Recognition ([1]), there is a divergent context-free language family to decide Clique as defined in Definition 3.2. Lemma 3.3 showed under the polynomial-time assumption that as n grows, there must be a state at level m in the minimal automaton for Clique formalized as in Equation 4 that has in-degree greater than 1. But by Lemma 2.1, there can be no state at level m for Clique formalized as Equation 4 with in-degree greater than 1. This contradiction blocks the existence of any polynomially bounded circuit for $\Gamma(n)$. Since Clique is an NP-Complete problem, this means that $P \neq NP$. \square

References

- [1] Raymond Greenlaw, H. James Hoover, and Walter L. Ruzzo. *Limits to Parallel Computation: P-Completeness Theory*. Oxford University Press, Oxford, 1995.
- [2] Michael Sipser. *Introduction to the Theory of Computation*. Cengage, Boston, third edition, 2013.